



جامعة القاهرة

LOGO.ADAM95.COM



---

# CO2 Emissions Data Analysis and Modeling

---

Name	ID
Rana Essam Ibrahim	20210133
Ahmed Yehia	20210049
Mohannad Hisham	20210413
NourEldin Ahmed	20210430
Merna Islam	20210500

**TA:** Eng. Sarah El-Nady

## a. Loading the Dataset

first, we import these libraries.

```
1 import pandas
2 import numpy as np
3 import matplotlib.pyplot as plot
4 from sklearn.linear_model import SGDClassifier
5 from sklearn.metrics import r2_score, accuracy_score
6 from sklearn.utils import shuffle
7 from sklearn.model_selection import train_test_split
8 from sklearn.preprocessing import LabelEncoder
9 import seaborn as sns
10 from sklearn.preprocessing import StandardScaler
```

To begin, we load the dataset and inspect its structure.

```
# load the data
data = pandas.read_csv('co2_emissions_data.csv')
data
```

✓ 0.0s

	Make	Model	Vehicle Class	Engine Size(L)	Cylinders	Transmission	Fuel Type	Consumption City (L/100 km)	Fuel Consumption Hwy (L/100 km)
0	ACURA	ILX	COMPACT	2.0	4	AS5	Z	9.9	6.7
1	ACURA	ILX	COMPACT	2.4	4	M6	Z	11.2	7.7
2	ACURA	ILX HYBRID	COMPACT	1.5	4	AV7	Z	6.0	5.8
3	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS6	Z	12.7	9.1
4	ACURA	RDX AWD	SUV - SMALL	3.5	6	AS6	Z	12.1	8.7
...	...	...	...	...	...	...	...	...	...
7380	VOLVO	XC40 T5 AWD	SUV - SMALL	2.0	4	AS8	Z	10.7	7.7
7381	VOLVO	XC60 T5 AWD	SUV - SMALL	2.0	4	AS8	Z	11.2	8.3
7382	VOLVO	XC60 T6 AWD	SUV - SMALL	2.0	4	AS8	Z	11.7	8.6
7383	VOLVO	XC90 T5 AWD	SUV - STANDARD	2.0	4	AS8	Z	11.2	8.3
7384	VOLVO	XC90 T6 AWD	SUV - STANDARD	2.0	4	AS8	Z	12.2	8.7

7385 rows x 13 columns

## b. Data Analysis

- i. check whether there are missing values
- ii. check whether numeric features have the same scale

```
1 # step 1 : data analysis
2
3 # check nulls in the data
4 print(data.isnull().sum())
5
6 # get description of the data to detect the scale of the data
7 numericFeatures = data.select_dtypes(include=[np.number])
8 print(numericFeatures.describe())
9
```

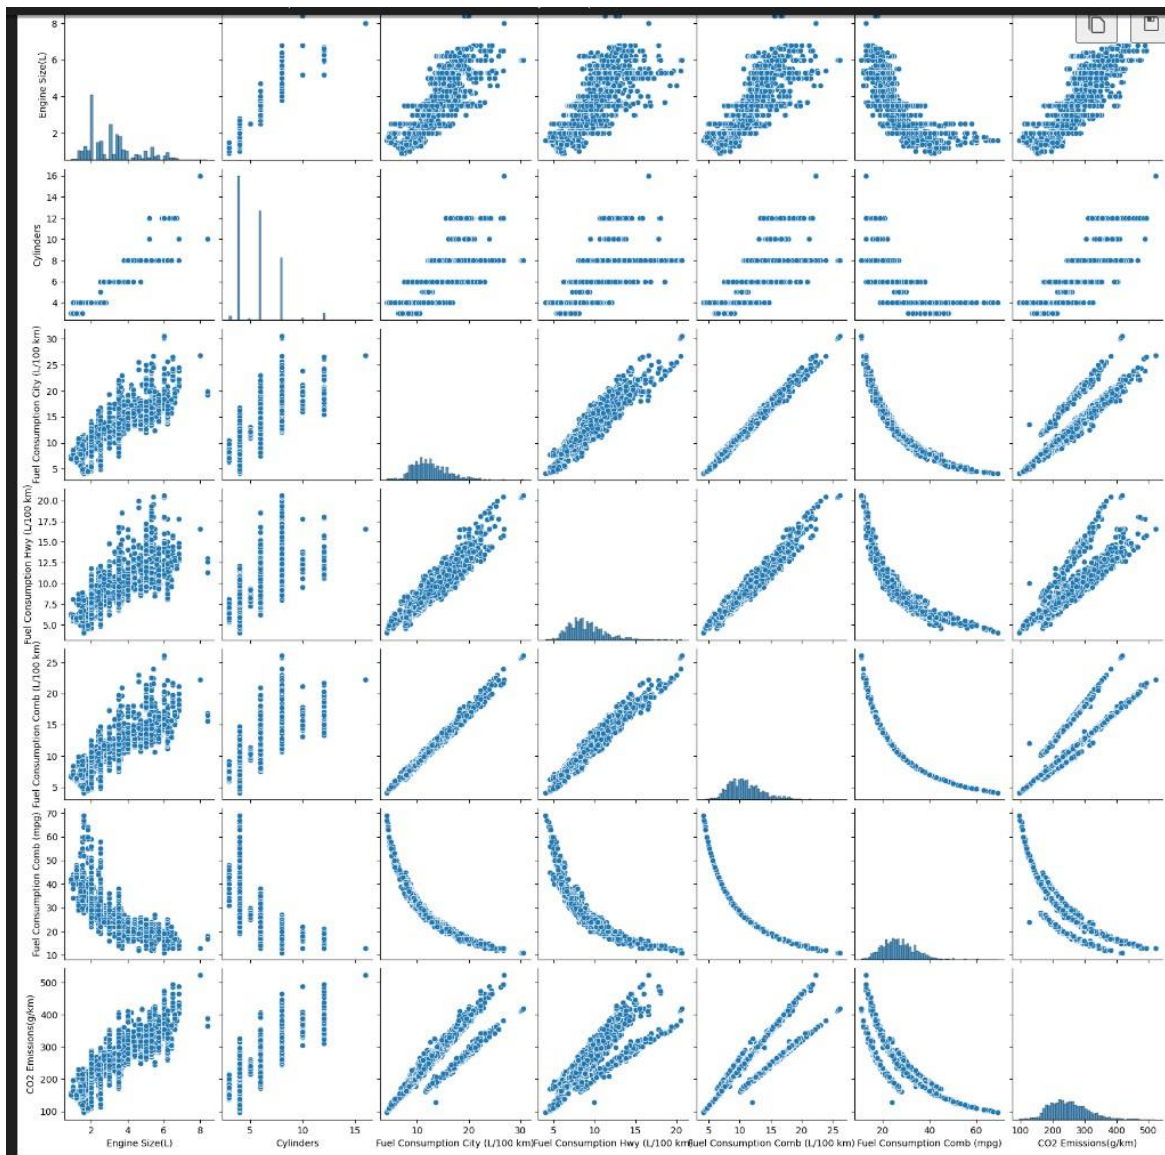
```
Make      0
Model     0
Vehicle Class  0
Engine Size(L)  0
Cylinders  0
Transmission  0
Fuel Type   0
Fuel Consumption City (L/100 km)  0
Fuel Consumption Hwy (L/100 km)  0
Fuel Consumption Comb (L/100 km)  0
Fuel Consumption Comb (mpg)  0
CO2 Emissions(g/km)  0
Emission Class  0
dtype: int64
Engine Size(L)  Cylinders  Fuel Consumption City (L/100 km)  \
count      7385.000000  7385.000000  7385.000000
mean         3.160068    5.615030    12.556534
std         1.354170    1.828307     3.500274
min         0.900000    3.000000     4.200000
25%         2.000000    4.000000    10.100000
50%         3.000000    6.000000    12.100000
75%         3.700000    6.000000    14.600000
max         8.400000   16.000000    30.600000

Fuel Consumption Hwy (L/100 km)  Fuel Consumption Comb (L/100 km)  \
...
25%                22.000000    208.000000
50%                27.000000    246.000000
75%                32.000000    288.000000
max                69.000000    522.000000
```

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...

### iii. Pairplot Visualization

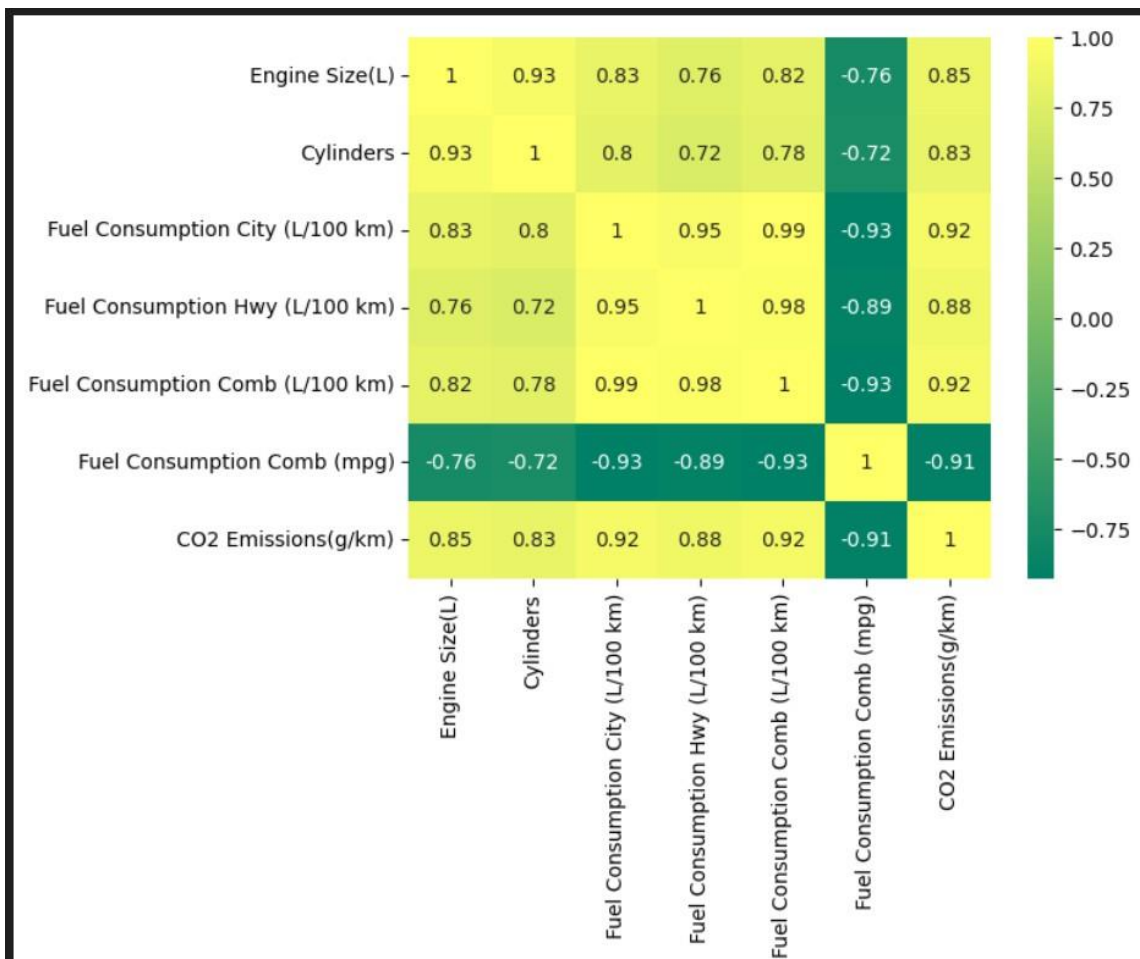
```
1 sns.pairplot(data, diag_kind='hist')
2 plot.show()
3
```



#### iv. Correlation Heatmap



```
1 # as the heatmap works numerical values so we well select the numerical values only
2 numericCols = data.select_dtypes(include=[np.number])
3 correlation_matrix = numericCols.corr()
4 sns.heatmap(correlation_matrix, annot=True, cmap='summer')
5 plot.show()
```



## c.Data Preprocessing

First, Select the features that we work with them and any variable we will use in any model

```
1 # this is the columns that we will use in the linear regression model and the logistic regression  
  model as independent variables  
2 feature1 = 'Cylinders'  
3 feature2 = 'Fuel Consumption Comb (L/100 km)'  
4 output = 'CO2 Emissions(g/km)' # which is used in linear regression as dependent variable  
5 alpha = 0.1  
6 maxIterations = 1000  
7 output2 = 'Emission Class' # which is used in logistic regression as dependent variable
```

- i. Target and Features separation has been done before training and testing for each model to be more readable as we are working on 2 models, not only one
- iii. Shuffle and split the data into training and test data

```
1 data = shuffle(data, random_state=100) # shuffle the data  
2  
3 trainingData, testData = train_test_split(data, test_size=0.3, random_state=100) # split the data  
  into training and testing data
```

iv. Scale the numerical values that we will use using Standard Scaler

```
1 scaler = StandardScaler()
2
3 # fit and transform the training data
4 trainingDataScaled = pandas.DataFrame(scaler.fit_transform(trainingData[[feature1, feature2, output]]), columns=[feature1, feature2, output])
5
6 # transform only the test data
7 testDataScaled = pandas.DataFrame(scaler.transform(testData[[feature1, feature2, output]]), columns=[feature1, feature2, output])
```

ii. Encode the categorical values (this step could be done before splitting) but it will not affect the result in both cases

```
1 # label encoder to encode the emission class
2 labelEncoder = LabelEncoder()
3
4 # encode the output2 values to be numerical
5 output2EncodedTrain = pandas.DataFrame(labelEncoder.fit_transform(trainingData[output2]), columns=[output2])
6 output2EncodedTest = pandas.DataFrame(labelEncoder.fit_transform(testData[output2]), columns=[output2])
7
8
9
```

Finally, concatenate both categorical values and numerical values after Preprocessing

```
1 # concatenat both numerical values with categorical values
2 trainingData = pandas.concat([trainingDataScaled, output2EncodedTrain], axis=1)
3 testData = pandas.concat([testDataScaled, output2EncodedTest], axis=1)
4
```

## d. Linear Regression Implementation and Training

Select the Values that will enter the model




```
1 feature1 = 'Cylinders'
2 feature2 = 'Fuel Consumption Comb (L/100 km)'
3 output = 'CO2 Emissions(g/km)' # which is used in linear regres
  sion as dependent variable
4 alpha = 0.1
5 maxIterations = 1000
```



```
1 x1 = trainingData[feature1].values
2 x2 = trainingData[feature2].values
3
4 x = np.column_stack((x1, x2))
5 y = trainingData[output].values
```




Create array of costs to store the cost after each iteration



```
1 costs = [] # to store the cost of each iteration
```

Implement the GD for linear regression model



```
1 def fitGD(x, y, alpha, maxIterations):
2     x = np.c_[np.ones(x.shape[0]), x] # Add a column of ones to x for the bias term
3     thetas = np.random.rand(x.shape[1])
4     for i in range(maxIterations):
5         h = np.dot(x, thetas) # the hypothesis function
6         for j in range(len(thetas)):
7             partialDerivative = (1/len(y)) * np.sum((h - y) * x[:, j]) # the partial derivative
of the cost function
8             thetas[j] = thetas[j] - alpha * partialDerivative # update the thetas values
9             cost = (1/len(y)) * np.sum(np.square(h - y)) # the cost function
10            costs.append(cost) # store the cost of each iteration
11
12     return thetas
13
```

Train the model to get thetas using training set




```
1  thetas = fitGD(x, y, alpha, maxIterations) # fit the model using the training data to get the thetas values
2  print(thetas)
3
```

```
[-1.22506363e-16  2.93912439e-01  6.88954267e-01]
```

Draw the cost function

```
1  
2 # draw the cost function  
3 plot.plot(costs)  
4 plot.xlabel('Number of Iterations')  
5 plot.ylabel('Cost')  
6 plot.show()  
7  
8  
9
```



## Evaluate the Linear Regression Model

```
1  #4.2 Predict the CO2 emissions for the test set using the linear regression model
2
3  def predictUsingLinearRegression(x, thetas):
4      return np.dot( x, thetas)
5
6  # Prepare the test data
7  x1Test = testData[feature1].values
8  x2Test = testData[feature2].values
9
10
11  xTestTemp = np.column_stack((x1Test, x2Test)) # 2 diemtional X
12  xTest = np.c_[np.ones(xTestTemp.shape[0]), xTestTemp] # Add a column of ones to x_test for the bias term
13
14  # Predict the CO2 emissions for the test set
15  yTest = testData[output].values
16
17  yPred = predictUsingLinearRegression(xTest, thetas)
18
19  # Calculate the R2 score
20  r2 = r2_score(yTest, yPred)
21  print('      "R2 score For"')
22
23  print(f'Linear Regression model: {r2}')
```

```
      "R2 score For"
Linear Regression model: 0.881053510923265
```

## e. Train the logistic Regression Model using Stochastic Gradient Descent

```
1 # Step 3.2 - Training the logistic regression model
2
3 # Predict the co2 emission class
4 yTest2 = testData[output2].values
5
6 # Stochastic Gradient Descent Classifier
7 logisticModel = SGDClassifier(loss='log_loss', random_state=99, max_iter=10000, tol=1e-3)
8
9 # Train the model using the training data
10 logisticModel.fit(x, trainingData[output2])
11
```

Evaluate the logistic regression model

```
1 # Predict the CO2 emission class for the test set
2 yPred2 = logisticModel.predict(xTestTemp)
3
4
5 accuracyScore = accuracy_score(yTest2, yPred2)
6 print(' "accuracy_score For"')
7 print(f'Logistic Regression model: {accuracyScore}')
8
9
```

```
• "accuracy_score For"
  Logistic Regression model: 0.9733754512635379
```